

Forward Deployed *Dossier*

Every role on the Notion job list, in full — what was on file, plus what the web says as of today. Each record separates the two: **On file** is the Notion entry verbatim in substance; **Update** is this pass’s research, including four material changes that affect whether a role still belongs in the pool.

APPROVED SLOTS	RANKED POOL	WATCH ITEMS	MATERIAL CHANGES
2 / 5	10	2	4

Approved by Chris

2 OF 5 SLOTS · GATE OPEN

The five-slot approved list. Nothing reaches this tier without an explicit judgment call in a voice session.

APPROVED

AD

Forward Deployed AI Engineer — Claude 2026

AIMPOINT DIGITAL

CONFIRMED LIVE

US REMOTE

SENIOR / L5

FULL-TIME

COMP NOT POSTED

LOCATION

United States, fully remote

COMPENSATION

Not stated in posting

REQUISITION

breezy 6845c658955b

APPLY AT

aimpoint-digital.breezy.hr

ON FILE — NOTION

US remote, full-time, senior L5. Helps enterprise clients move from strategy into production-scale Claude and agentic AI systems — architect, prototype, build, evaluate, operationalize. Explicitly wants human-in-the-loop workflows, approval chains, and agent supervision patterns. Strong fit on the AI-tooling axis (Claude, agentic) and the prototype-to-production axis; smaller consultancy brand. Approved by Chris.

UPDATE – WEB, 30 AUG 2026

Still live on Breezy and mirrored across Jobgether, fwddeploy.com, RemoteRocketship and Himalayas. Aimpoint describes itself as a data, AI, analytics and operations-research advisory and solution-engineering firm building enterprise-grade data and AI platforms, decision-intelligence systems and production AI applications.

The scope reads broader than the Notion note captured: alongside client delivery, the role builds **reusable solution patterns, accelerators, technical demos, reference architectures and thought leadership for Aimpoint’s Anthropic practice**, and helps clients define an AI operating model covering governance, model evaluation and adoption workflows. Requirements confirm the HITL emphasis and add AI risk, reliability, security and compliance fluency plus production software-engineering fundamentals – APIs, data pipelines, enterprise systems.

Second entry point at the same firm: **Lead Forward Deployed Engineer, Databricks 2026 (US/UK)**, breezy fa7df72d90d8 – a Lead-level title, worth a look if the Claude req stalls.

APPROVED



Forward Deployed Engineer — Agentforce

SALESFORCE

CONFIRMED LIVE NOW A FAMILY OF REQS SF PASSES FILTER TRAVEL 25-50%

LOCATION	BASE SALARY	REQUISITIONS
Varies by req — San Francisco passes the filter; Seattle / NYC variants do not	\$180,200–\$247,900 select SF & NYC metro cities \$150,100–\$227,000 elsewhere	JR349466 · JR339744 · JR343861
APPLY AT careers.salesforce.com		

ON FILE – NOTION

Recorded as approved on the *Approved job shortlist* gate — “2 of 5 roles approved (Salesforce, Aimpoint Digital).” The candidate-pool note ranked it among the two strongest on the rubric for **brand + agentic AI + remote options**. The pool page no longer carries a full candidate write-up for it, so this record is rebuilt from the gate entry plus today’s research.

UPDATE — WEB, 30 AUG 2026

This is not one posting but a **family of Agentforce FDE requisitions**, which changes how to approach it — pick the req, not the company:

JR349466 Forward Deployed Engineer, all levels (mid/senior) · JR339744 FDE, Agentforce for Supply Chain — San Francisco, Seattle, Chicago, Austin, Atlanta, New York, Indianapolis · FDE, Agentforce Orchestration — New York and Seattle · JR343861 FDE, Australia. KPMG also hires a Salesforce Health Cloud & Agentforce FDE through the partner channel.

The work: designing, implementing and optimizing automated business-process solutions — Salesforce calls them “Blueprints” — on the Agentforce platform, deploying enterprise-grade agentic AI inside customer environments, and owning the roadmap and feature lifecycle for scalable process solutions. Requirements: hands-on LLM integration, AI orchestration frameworks and prompt engineering; data modeling and integration on Data Cloud, Snowflake or Databricks; **direct Salesforce platform configuration, customization and development**; travel 25–50% to customer sites.

Open question: the Salesforce-platform requirement is the one axis with no track record behind it. Everything else on this req is a clean match; that line is the screen to prepare for.

NEXT ACTION

Target JR339744 on the San Francisco location — it is the only variant in the family that clears the LA / Santa Monica / Tempe / SF / remote filter.

Ranked pool

10 UNREVIEWED CANDIDATES

The wide discovery layer, ranked loosely, reviewed top-down. Nothing here is vetted.

Location filter is mechanical: Los Angeles, Santa Monica, Tempe, San Francisco, or remote — anything else drops automatically.



Forward Deployed Engineer, Applied AI

ANTHROPIC

CONFIRMED LIVE

COMP ESTIMATE REVISED DOWN

SAN FRANCISCO & NYC

TRAVEL 25–50%

4+ YRS CUSTOMER-FACING

LOCATION

San Francisco, CA — also listed
New York

COMPENSATION

Not posted. Third-party 2026:
\$350K–\$550K TC mid–senior
FDE

REQUISITIONS

gh 5391021008 • 5391016008
• 5302966008 • 5074695008
(Custom Agents) •
5079562008 (Federal
Civilian) • 5012991008

APPLY AT

anthropic.com/careers/jobs

ON FILE – NOTION

Embeds directly with strategic enterprise customers to drive AI adoption: builds production applications with Claude inside customer systems and ships technical artifacts — **MCP servers, sub-agents and agent skills** — into production workflows. Requires 4+ years in a technical customer-facing role. Travel 25–50% to customer sites.

Fit note on file: the single closest match in the pool to the Career Portfolio rubric — customer-facing, prototype-to-production, explicitly agentic — running on the exact stack Chris already uses hands-on daily (Claude Code, MCP, sub-agents, agent orchestration). The travel is a feature, not a cost, against his stated goal of work with real travel built in. Comp recorded as roughly \$665K–\$750K for FDE-track L4–L5, with a wider Applied AI band of \$300K–\$1.2M — flagged at the time as unverified third-party estimate.

UPDATE – WEB, 30 AUG 2026

Two further live Greenhouse reqs surfaced beyond the five on file: 5391016008 and 5012991008. The role is syndicated across VC job boards — Accel 67504565, General Catalyst and Menlo Ventures 69674588 — so the funnel is wider than the careers page alone suggests. Anthropic brands its FDEs internally as “**Applied AI Engineers**,” which is the search term that surfaces the most listings.

Role description confirms the Notion entry: embed with strategic customers to drive transformational AI adoption, ship advanced AI applications inside customer systems, combining engineering depth, frontier-AI understanding and customer-facing skill while holding safety and reliability standards.

Compensation correction: current sources put mid-to-senior FDE total comp at \$350K–\$550K, with senior FDEs reaching north of \$785K. A posted Anthropic Pre-Sales Program Lead, Forward Deployed Engineering in SF lists \$270K–\$365K base. levels.fyi shows Anthropic SWE in the SF Bay Area at \$555K–\$767K+, but that is the software-engineering ladder, not the FDE track. The \$665K–\$750K figure on file looks high for L4–L5; plan against \$350K–\$550K.

Sr. Forward Deployed Engineer — Communications, Media, Entertainment & Games

DATABRICKS

CONFIRMED LIVE

COMP CONFIRMED

REPOSTED UNDER NEW REQ IDS

LOS ANGELES, CA

LOCATION

Los Angeles, CA 90013 — on
the target list

BASE SALARY

\$182,000–\$250,208 + annual
bonus, equity, benefits

REQUISITIONS

gh_jid 8461168002 ·
8341311002 · (on file:
8432827002)

ALSO LISTED

ziprecruiter
7252996b2380ddb7 · indeed
jk 8ef34cdf735af006

ON FILE – NOTION

Works side by side with senior client technical and functional staff to **rapidly prototype and deliver high-impact GenAI-enabled solutions**, owning architecture and end-to-end design across data engineering, AI and application development. Base pay \$182,000–\$250,208 across zones, plus annual bonus, equity and benefits. One LA senior FDE req noted a recruiting close of September 30, 2026.

Fit note on file: clean rubric match on all three axes, in a target city, with an explicit prototype-to-production mandate. The media and entertainment vertical is a domain with no track record behind it, which is the main open question.

UPDATE – WEB, 30 AUG 2026

Confirmed live and **currently posted under two requisition IDs neither of which is the one on file**: 8461168002 and 8341311002, both under Professional Services & Operations on databricks.com/company/careers. The recorded 8432827002 no longer surfaces — Databricks rotates these reqs the same way AWS rotates PACE, so treat the role type as durable and the ID as disposable.

Pay range confirmed verbatim at **\$182,000–\$250,208 USD**, described as the expected base for non-commissionable roles, with eligibility for annual performance bonus, equity and benefits on top. Actual packages vary by job-related skills, depth of experience, relevant certifications and training, and specific work location — LA sits in the upper zone.

NEXT ACTION

Apply against 8461168002, the freshest of the two. If the Sept 30 close on file is accurate, this is the only pool entry with a known deadline — roughly four weeks out.

Forward Deployed Engineer (Enterprise AI Solutions Architect)

RESILINC

NOW INDIA-REMOTE, NOT US

SUBSTITUTE REQ IDENTIFIED

6+ YRS REQUIRED

FULLY REMOTE COMPANY

LOCATION ON FILE

United States, remote

LOCATION TODAY

India, remote — for this exact title

REQUISITIONS

lever 3027ee22... (India) ·
lever 1de72f37... (US-eligible substitute)

COMPENSATION

Not posted

ON FILE — NOTION

US remote. Sits at the intersection of customer delivery, product discovery, data engineering, applied AI and software engineering — exactly the customer-facing prototype-to-production shape. 6+ years required; stands out with agentic systems, LangChain/LangGraph and LLM-powered apps. Fully remote with in-person connect opportunities. Strong match on the rubric, weaker on brand — not a top-ten name.

UPDATE — WEB, 30 AUG 2026

Material change. The live Lever requisition for this exact title is now scoped “Senior Forward Deployed Engineer (Enterprise AI Solutions) — India (Remote),” req 3027ee22-5c12-4dfc-9cbd-bb353ada3cb9, syndicated the same way on Built In, fwddeploy and the Partech portfolio board. No US-scoped version of this title surfaced this pass. As recorded, the entry no longer clears the location filter.

The nearest same-shape opening that remains reachable is **Enterprise Solutions Architect — Agentic AI Deployment Services**, lever 1de72f37-2e8e-434a-ac0e-9add6785300f.

Role content otherwise matches the file: deliver complex enterprise deployments, working with strategic customers and cross-functional teams across Agent Success, Product, Engineering, Data, Support and Sales on problems past standard implementation — understand the problem, define the technical path, build what is needed, deploy to production, and leave behind something maintainable, documented, observable and supportable.

Company context: agentic AI for supply-chain risk prediction and real-time response, named a 2025 Gartner Magic Quadrant Leader, customers in life sciences and pharma, aerospace and defense, high tech and automotive; fully remote organization.

NEXT ACTION

Either swap this entry for lever 1de72f37 after confirming US eligibility, or drop it from the pool.
Do not spend effort on the India req.



Senior Applied AI Solutions Architect — Amazon Connect

AMAZON / AWS

DEADLINE PASSED 26 JUN 2026

FRESHER SIBLINGS OPEN

SF & MOUNTAIN VIEW, CA

AI VELOCITY TEAM

LOCATION

San Francisco and Mountain View, CA

BASE SALARY

\$176,000–\$239,000

REQUISITION ON FILE

10435917

LIVE SIBLINGS

10459921 · 10446686 ·
10393427 · 10479997 (Public Sector)

ON FILE – NOTION

Job ID 10435917. San Francisco and Mountain View, CA. Base \$176k–\$239k. Embedded with customers building agentic AI agents on Amazon Connect: model selection, prompt engineering, MCP servers, multi-agent orchestration, production deployment. Pure hands-on, AI Velocity Team.

UPDATE – WEB, 30 AUG 2026

The req is still visible on amazon.jobs as “Senior Applied AI Solutions Architect — Amazon Connect, Applied AI SA -AIVT” and on Glassdoor for San Francisco, but it **carries an application deadline of 26 June 2026** — two months past. Treat the ID on file as stale even though the page still renders.

Organizational detail the file did not have: the role sits in the **AWS Specialist and Partner Organization (ASP)**, where specialists own end-to-end go-to-market for their technology domain. The AI Velocity Team is described as a service-specific model that assigns dedicated advisory and hands-on development resources directly to customers to reach production-ready outcomes in weeks rather than months — the same compress-the-timeline framing as PACE, aimed at Connect.

Four fresher reqs of the same role are open: 10459921, 10446686 and 10393427 (Sr. Applied AI SA, Amazon Connect) plus 10479997, a Public Sector variant.

NEXT ACTION

Retire 10435917 from the pool and re-point at 10459921 or 10446686 after confirming a CA location on each.



Sr. GenAI Specialist Solutions Architect, Agentic AI

AMAZON / AWS

CONFIRMED LIVE

REQ RESCOPED TO WORLDWIDE GTM

SF & SANTA CLARA, CA

KIRO – DIRECT EVIDENCE MATCH

LOCATION	BASE SALARY	REQUISITION
San Francisco and Santa Clara, CA	\$176,000–\$239,000	10431946

RELATED OPEN REQs
10425360 (Agentic Engineering / Kiro) · 10445643 · 10470183

ON FILE – NOTION

Job ID 10431946. San Francisco and Santa Clara, CA. Base \$176k–\$239k. Customer-facing enablement on agentic coding tools like Kiro: demos, hackathons, reference implementations, turning field feedback into product requirements.

Confirmed live on amazon.jobs, including a Japanese-locale listing under the same ID. Note that 10431946 now also surfaces re-titled as **“Worldwide Specialist Solutions Architect — Agentic Development, Data & AI GTM,”** so the req has been rescoped toward worldwide GTM since it was captured. Confirm the posted location before applying — a worldwide-GTM rescope is exactly the kind of change that moves a role to Seattle.

Current description: act as subject-matter expert for customers designing solutions on agentic AI services to transform their software development lifecycle; work with customers directly to understand their business problems and SDLC; implement generative AI solutions using agentic coding tools like Kiro; deliver briefings, demos, hackathons and deep-dive sessions.

Three adjacent reqs are open in the same space: 10425360 Senior Specialist SA, Agentic Engineering / Kiro — the closest pure-Kiro role — plus 10445643 Sr AI Solution Architect and 10470183 Sr. GenAI Specialist PSA.

Strongest evidence match in the Amazon group: Kiro is already in daily hands-on use, so the “experience with AI-driven development tools” preference is answered with current work rather than a claim.

6

Sr. Prototyping Architect, PACE

AMAZON / AWS

CONFIRMED LIVE

OLD TEAM NAME — OLDER REQ

THE DURABLE TARGET

LOCATION UNCONFIRMED

REQUISITION

3169823

TEAM

Listed as Prototyping *and*
Cloud Engineering — the pre-
rename name

COMPENSATION

Not posted on this req

CURRENT SIBLINGS

10414556 • 10449045 •
10470048 • 3206379 •
10410797 (TPM)

Job ID 3169823. Fresh repost of the standard PACE role, same shape as the durable target documented on the Career Portfolio page. That page names PACE Prototyping Architect and Senior Prototyping Architect — originally job ID 10470042 — as the top-ranked target as of late August 2026, with the standing instruction to treat **the role type as the durable target, not any single job ID**, since PACE reposts constantly (3169823, 3206379 and 10414556 have all appeared).

What the role does, per the portfolio page: pairs architects directly with AWS customers to understand strategic business goals and then rapidly build working prototypes — machine learning, generative AI, autonomous agents and multi-agent orchestration, RAG, IoT, serverless. Those prototypes become the foundation for new customer products, drive service adoption, and feed customer needs back into AWS service-team roadmaps. Why it fits: explicitly greenfield rather than the brownfield operational work of the last two years, and it leans on the teaching background, since much of the role is helping customers understand the art of the possible.

UPDATE — WEB, 30 AUG 2026

Still listed, but under the title “**Sr. Prototyping Architect, PACE, AWS Prototyping and Cloud Engineering**” — the pre-rename team name. That makes 3169823 an older req than the renamed ones, not a fresh repost as recorded. 3206379 carries the same old name.

The PACE description itself has moved decisively toward agents. Current reqs open with: the team “**transforms ambitious ideas into working Generative AI and Agentic AI prototypes in days and weeks, not months,**” and describe hands-on work architecting production-ready AI prototypes using prompt engineering, RAG architectures, function calling, autonomous agents and complex multi-agent orchestration on AWS AI/ML services. The profile asked for is a hands-on full-stack builder with deep cloud expertise, comfortable across multiple languages and frameworks in emerging areas including AI/ML, edge compute, robotics and spatial computing.

Freshest US-shaped sibling to chase: 10414556, Senior Prototyping Architect, Prototyping and AI Customer Engineering. A 10470048 Physical AI variant is also open.



Prototyping Architect, PACE

AMAZON / AWS

HIRING ENTITY IS BRAZIL

FAILS LOCATION FILTER

RECOMMEND DROP

LOCATION ON FILE

Not recorded

LOCATION TODAY

**Amazon AWS Services Brazil
Ltda.**

REQUISITION

10379831

SUGGESTED REPLACEMENT

10414556 or 10470048

ON FILE — NOTION

Job ID 10379831. Another current PACE repost. Explicitly mentions Agentic AI in the description — flagged as worth noting since that is a growing emphasis in the role.

UPDATE — WEB, 30 AUG 2026

Material change. The req is posted under Amazon AWS Services Brazil Ltda. The Agentic AI emphasis recorded on file is real — the description covers prototypes built with agentic AI, machine learning, deep learning, generative AI, IoT and serverless architectures, on an agile team that partners with customers to experiment with new ideas — but the hiring entity fails the LA / Santa Monica / Tempe / SF / remote filter outright.

NEXT ACTION

Drop from the pool on the mechanical filter. Replace with 10414556 or 10470048 once a US location is confirmed on either.

8



Forward Deployed Engineer, Professional Services

CLOUDFLARE

CONFIRMED LIVE

STAFF AUGMENTATION, NOT PROTOTYPING

NO AI CONTENT IN REQ

DISTRIBUTED / REMOTE

LOCATION

Distributed — a hybrid variant
also appears on some boards

COMPENSATION

Not posted

REQUISITION

greenhouse 8050386

RUBRIC VERDICT

Weakest match in the pool

ON FILE — NOTION

Cloudflare, distributed/remote. Same rubric as PACE per the Career Portfolio page: customer-facing, prototype-to-production. **Needs a deeper look at seniority and comp.**

The open question on file is now answered, and the answer is unfavorable. The req is live and actively hiring, but Cloudflare frames the FDE program explicitly as **“not a traditional support or purely advisory engagement — it is a staff-augmentation model”**: a dedicated hands-on expert embedded in the customer’s environment, acting as an extension of their team, managing, optimizing and evolving *their existing Cloudflare deployment*.

The deliverables named are Solution Designs, HLDs, LLDs, databuilds, procedures, test plans, deployment plans, migration plans and as-builts; leading technical workshops, knowledge-transfer sessions and special-event prep for product launches, traffic events, security drills and quarterly business reviews; and acting as voice of the customer into Engineering, Support and Product. Seniority is genuinely high — recognized subject-matter expert, owns the most complex high-stakes engagements, architectural decisions on mission-critical infrastructure, executive stakeholders, high autonomy.

But this is **delivery-and-operate work on an existing platform, not greenfield prototype-to-production**, and the requisition carries no AI or agentic content at all. It fails the rubric on the axis that matters most, despite the matching job title.

NEXT ACTION

Recommend dropping to the bottom of the pool or removing it. The title matches; the work does not.



Agent Engineer

SIERRA

CONFIRMED LIVE

BASE BAND WIDER AND LOWER THAN FILED

SF, PRIMARILY ONSITE

EQUITY AT \$15.8B VALUATION

LOCATION

San Francisco, CA — primarily onsite, some flexibility for occasional remote

BASE SALARY

\$150K–\$300K posted (file recorded \$225K–\$300K)

TOTAL COMP

levels.fyi: \$200K–\$520K+ TC, median around \$460K

REQUISITIONS

ashby 76617113-a408-40ec-9c77-aa1174454a37 · e9f5fdb6-91ee-4c55-9230-41ec8865650e

Sierra’s agent engineers are FDE-shaped: hands-on customer deployment of conversational AI agents into production. Two Ashby listings recorded — Agent Engineer, and Software Engineer, Agent (Healthcare). Reported base roughly \$225K–\$300K with the remainder in options tied to Sierra’s valuation; noted as a third-party figure, not a posted range.

Fit note on file: strongest agentic-AI signal outside the frontier labs and a genuine builder role, but the weakest of the three on seniority framing — titled as an IC engineer rather than architect or senior FDE — and onsite SF is the least flexible location in the pool.

UPDATE – WEB, 30 AUG 2026

Ashby req e9f5fdb6 is still live. One posting shows a base range of **\$150K–\$300K** — wider, and considerably lower at the floor, than the \$225K–\$300K on file. Treat the bottom of that band as real: at an IC engineer title, an offer can land well below the filed figure.

Total compensation tells a better story: levels.fyi puts Sierra software engineers in the SF Bay Area at roughly \$200K entry to \$460K+ at senior, exceeding \$520K at the top, with a median near \$460K — and 100% of Sierra postings reference equity, which sits outside those medians.

Sierra’s valuation is now reported at \$15.8B (Bret Taylor’s company), which is the number any option grant is struck against — the single most important variable in this offer. Benefits listed: 401(k) with match, generous parental leave, flexible personal stipend. Listings run through sierra.ai/careers and the Sequoia job board.



Sr. Worldwide GTM Specialist,
Agentic AI — Kiro & Frontier Agents

AMAZON / AWS

SEATTLE ONLY – FAILS FILTER

REPOST WATCH

SUBSTITUTE IDENTIFIED

LOCATION	REQUISITION	COMPENSATION
Seattle only — no CA, AZ or remote option	10482606	Not posted
SUBSTITUTE		
10425360 – Senior Specialist SA, Agentic Engineering / Kiro		

Job ID 10482606. Seattle only — does not meet the location filter, so flagged for repost watch only, not added to the active pool. Kept in the numbering as the tenth entry.

UPDATE — WEB, 30 AUG 2026

No CA, AZ or remote variant surfaced this pass, so the exclusion stands. The nearest location-passing equivalent on the same subject matter is **10425360, Senior Specialist Solutions Architect, Agentic Engineering / Kiro** — the same Kiro and frontier-agents territory, sitting in the Solutions Architecture organization rather than Worldwide GTM, which is arguably the better org for hands-on work anyway.

NEXT ACTION

Keep 10482606 on repost watch; pursue 10425360 in the meantime.

Watch item

NOT IN THE POOL

Held outside the numbered pool because the location could not be confirmed when it was captured. It can now be resolved.

WATCH



Sr. Prototyping Architect, AWS Prototyping and AI Customer Engineering (PACE)

AMAZON / AWS

LOCATION NOW RESOLVED

CANADA ENTITY — FAILS FILTER

TEAM RENAME CONFIRMED

LOCATION ON FILE

Could not be confirmed

LOCATION TODAY

**Amazon Web Services
Canada, Inc.**

REQUISITION

10449045

RENAME EVIDENCE

10449045 • 10414556 •
10470048 • 10410797

ON FILE — NOTION

Not added to the pool: the mechanical filter requires a confirmed LA / Santa Monica / Tempe / SF / remote location, and that pass could not confirm the posted location for this req. Flagged anyway because **the team name has changed** — PACE now reads “Prototyping and *AI Customer Engineering*” rather than “Prototyping and Cloud Engineering.” That is a live signal about where the top durable target is heading, and it should be reflected on the Career Portfolio page once confirmed. Instruction left for the next agent: confirm location on 10449045 and add it if it passes the filter.

Both open questions are now answered.

Location: the req is posted under **Amazon Web Services Canada, Inc.** It fails the filter, the same way 10379831 does on Brazil. Mark it excluded by location rather than unconfirmed, and close the instruction left for the next agent.

Rename: confirmed across four current requisitions — 10449045, 10414556, 10470048 and TPM 10410797 all carry “AWS Prototyping and AI Customer Engineering.” The older 3169823 and 3206379 still carry the Cloud Engineering name, which dates them. **The Career Portfolio update is now safe to make.**

NEXT ACTION

Update the Career Portfolio page: PACE = AWS Prototyping *and* AI Customer Engineering. Then chase 10414556 as the US-location PACE candidate.

Standing issues

FOR THE NEXT JUDGMENT SESSION

What this pass changes

- **Four material changes.** Resilinc moved to India-remote; PACE 10379831 is a Brazil entity; PACE 10449045 is a Canada entity; Amazon Connect 10435917 passed its 26 June application deadline. Four of the ten pool entries are no longer applicable as recorded — every one has a named substitute above.
- **Two estimates corrected.** Anthropic FDE total comp reads \$350K–\$550K mid-senior on current sources, not the \$665K–\$750K on file. Sierra’s posted base floor is \$150K, not the \$225K on file.
- **The Amazon concentration question resolves itself.** The Critic’s “at most two per employer family” rule was flagged as a Should-fix with four Amazon entries in the pool. With 10379831 failing on Brazil, 10482606 failing on Seattle and 10435917 past deadline, the natural survivors are 10431946 (Agentic AI / Kiro) and one PACE req, most likely 10414556 pending location. That lands on two, without a judgment call — but the call is still Chris’s to confirm.
- **Cloudflare fails on substance, not location.** The “deeper look at seniority and comp” noted on file turned up a staff-augmentation model with no AI content. It is the one entry that matches the rubric on title alone.
- **Approved gate still open at 2 of 5.** On this pass’s evidence the strongest unapproved candidates are Anthropic FDE (closest rubric match, exact stack) and Databricks LA (only entry in a target city with a confirmed range and a known deadline).

Verification level

- Every record above was cross-checked by web search on 30 August 2026. As on the 29 August pass, **this session's network proxy blocks direct fetches of amazon.jobs and job-boards.greenhouse.io**, so no listing was confirmed live at the URL itself. Treat verification as search-level, not URL-level — reliable for existence, titles, locations and posted ranges, but confirm deadlines and exact locations on the listing page before applying.
- Compensation figures marked as third-party (Anthropic, Sierra total comp) come from levels.fyi and aggregator reporting, not employer postings. Posted ranges — Databricks, Salesforce, the two Amazon specialist roles — are employer-stated.

Sources consulted

- Notion — Job Candidate Pool (Unreviewed)
- Notion — Approved job shortlist gate
- Notion — Career Portfolio
- Anthropic — Forward Deployed Engineer (Greenhouse)
- Anthropic FDE, Applied AI (Accel board)
- levels.fyi — Anthropic, SF Bay Area
- Databricks — Sr. FDE, CME&G (8461168002)
- Databricks LA listing (ZipRecruiter)
- Aimpoint Digital — FDE AI Engineer, Claude 2026
- Aimpoint Digital — Lead FDE, Databricks 2026
- Salesforce — FDE, Agentforce for Supply Chain
- Salesforce — FDE, all levels (JR349466)
- AWS — Sr. Applied AI SA, Amazon Connect (10435917)
- AWS — Sr. GenAI Specialist SA, Agentic AI (10431946)
- AWS — Senior Specialist SA, Agentic Engineering / Kiro (10425360)
- AWS — Sr. Prototyping Architect, PACE (3169823)
- AWS — Prototyping Architect, PACE (10379831)
- AWS — Sr. Prototyping Architect, PACE (10449045)
- AWS — Senior Prototyping Architect, PACE (10414556)

- Cloudflare — FDE, Professional Services (8050386)
- Sierra — Software Engineer, Agent (Ashby)
- levels.fyi — Sierra, SF Bay Area
- Resilinc — Senior FDE, Enterprise AI Solutions (India)
- Resilinc — Enterprise Solutions Architect, Agentic AI Deployment

Compiled 30 August 2026 from the Notion Job Candidate Pool, the Approved job shortlist gate and the Career Portfolio rubric, with a same-day web pass over every listing. On file text preserves the substance of the Notion entries so the record stays useful after a listing closes; Update text is this pass's research and is dated accordingly. Requisition IDs rotate — PACE and Databricks both repost the same role under new IDs regularly — so treat the role type as the durable target and the ID as disposable.